

# Multilayer Perceptron in Verilog with Fixed-Point Backpropagation(Task16)

Adarsh Tiwari

April 20, 2025

## Abstract

This report presents a minimal 2-2-1 multilayer perceptron (MLP) implemented in Verilog, trained on a small dataset via fixed-point backpropagation and tested on one held-out sample. The network architecture has been described, mathematical foundations of gradient descent, the use of an 8.8 fixed-point sigmoid lookup table, and key implementation details. The design omits bias terms and employs an identity activation approximation in hardware, illustrating how basic neural-network concepts map to RTL.

## 1 Introduction

Artificial neural networks (ANNs) are computing systems inspired by the biological neural networks of animal brains. A multilayer perceptron (MLP) is a type of feedforward ANN consisting of fully connected layers with nonlinear activations [1, 2]. In hardware, MLPs can be implemented using fixed-point arithmetic and lookup tables to approximate activation functions [5].

## 2 Network Architecture

Our network has:

- Two 16-bit inputs  $x_1, x_2$ .
- One hidden layer with two neurons  $h_1, h_2$ .
- One output neuron  $y$ .
- No bias terms.
- Sigmoid activations in hidden and output layers implemented via an 8-bit LUT (256 entries) in 8.8 fixed point [5].




Figure 1: Architecture of the 2-2-1 multilayer perceptron (inputs  $x_1, x_2$ ; hidden neurons  $h_1, h_2$ ; output  $y$ ).

## 2.1 Layer Computations

$$\begin{aligned} h_1 &= \sigma(w_{11}x_1 + w_{21}x_2), \\ h_2 &= \sigma(w_{12}x_1 + w_{22}x_2), \\ y &= \sigma(w_{31}h_1 + w_{32}h_2), \end{aligned}$$

where all multiplications and additions are in 8.8 fixed-point format and  $\sigma(\cdot)$  denotes the sigmoid LUT [5].

## 3 Mathematical Formulation

### 3.1 Loss Function

I have used mean-squared error (MSE):

$$E = \frac{1}{2}(y - y_{\text{target}})^2.$$

### 3.2 Gradient Descent Backpropagation

In backpropagation, I calculated gradients layer by layer using the chain rule. Below I have derived each term explicitly and present the light updates.

#### 3.2.1 Sigmoid and Loss Derivatives

The sigmoid function

$$\sigma(u) = \frac{1}{1 + e^{-u}}$$

and its derivative

$$\sigma'(u) = \sigma(u)(1 - \sigma(u))$$

are implemented via LUT [5]. For the MSE loss,

$$\frac{\partial E}{\partial y} = (y - y_{\text{target}}).$$

### 3.2.2 Output Layer Gradients

Let the net input to the output neuron be

$$u_{\text{out}} = w_{31}h_1 + w_{32}h_2.$$

Then

$$\delta_{\text{out}} = \frac{\partial E}{\partial \nu_{\text{out}}} = \frac{\partial E}{\partial y} \sigma'(\nu_{\text{out}}) = (y - y_{\text{target}}) y (1 - y).$$

The gradients with respect to the output-layer lights are

$$\frac{\partial E}{\partial w_{3i}} = \delta_{\text{out}} h_i, \quad i = 1, 2,$$

i.e.

$$\Delta w_{31} = -\eta (y - y_{\text{target}}) y (1 - y) h_1, \Delta w_{32} = -\eta (y - y_{\text{target}}) y (1 - y) h_2.$$

### 3.2.3 Hidden Layer Gradients

For each hidden neuron  $h_i$ , net input

$$\nu_{h_i} = w_{1i}x_1 + w_{2i}x_2,$$

it holds

$$\frac{\partial E}{\partial \nu_{h_i}} = \delta_{\text{out}} w_{3i} \sigma'(\nu_{h_i}) = (y - y_{\text{target}}) y (1 - y) w_{3i} h_i (1 - h_i).$$

Hence light gradients

$$\frac{\partial E}{\partial w_{ji}} = \delta_{h_i} x_j, \quad j = 1, 2, \quad i = 1, 2,$$

leading to

$$\Delta w_{11} = -\eta (y - y_{\text{target}}) y (1 - y) w_{31} h_1 (1 - h_1) x_1$$

,

$$\Delta w_{21} = -\eta (y - y_{\text{target}}) y (1 - y) w_{31} h_1 (1 - h_1) x_2$$

,

$$\Delta w_{12} = -\eta (y - y_{\text{target}}) y (1 - y) w_{32} h_2 (1 - h_2) x_1$$

,

$$\Delta w_{22} = -\eta (y - y_{\text{target}}) y (1 - y) w_{32} h_2 (1 - h_2) x_2.$$

These explicit formulas guide the fixed-point updates in hardware.

## 4 Fixed-Point Sigmoid Lookup Table

The sigmoid function

$$\sigma(u) = \frac{1}{1 + e^{-u}}$$

is stored in a 256-word LUT covering  $[-8, 8]$  in 8.8 format. Input  $u$  maps to index  $0 \dots 255$ ; outputs use linear interpolation for off-grid values [5].

## 5 Verilog Implementation Details

- **Fixed-Point Multiplier:**  
`fixed_point_mult(a,b)` multiplies two 16-bit 8.8 numbers and extracts bits [23:8] of the 32-bit product [5].
- **Light Registers:** Six 16-bit regs hold  $w_{11}, w_{12}, w_{21}, w_{22}, w_{31}, w_{32}$ .
- **Forward Pass:** Purely combinational in an `always @(*)` block for  $h_1, h_2, y$ .
- **Backprop/Update:** Synchronous in an `always @(posedge clk)` block when `train` is high and not done.
- **Training Control:** A 32-bit counter runs up to `NUM_TRAINING_ITERATIONS` (1000) then raises `training_done`.

## 6 Future Work

Three RTL optimization directions emerge from this baseline:

1. **Pipelined Gradient Computation:**  
Add 3-stage pipeline for forward/backward passes  
Estimated  $2.8 \times$  throughput improvement
2. **Light Quantization:**  
Explore 4.12 fixed-point format vs  
8.8 Mixed-precision multipliers (QKeras-style)
3. **On-Chip Learning Circuits:**  
Batch normalization hard macros  
Gradient checkpointing to reduce memory

## 7 Results

```
# KERNEL: Test Result:
# KERNEL: Inputs: x1 = 0.500000, x2 = 0.500000
# KERNEL: Expected output: 0.000000
# KERNEL: Actual output: 0.515625
# KERNEL: Final weights:
# KERNEL: w1_1 = 191.585938
# KERNEL: w1_2 = 191.601562
# KERNEL: w2_1 = 196.621094
# KERNEL: w2_2 = 196.062500
# KERNEL: w3_1 = 179.878906
# KERNEL: w3_2 = 179.941406
# RUNTIME: Info: RUNTIME_0068 testbench.sv (125): $finish called.
# KERNEL: Time: 3140 ns, Iteration: 0, Instance: /mlp_tb, Process: @INITIAL#39_1@.
# KERNEL: stopped at time: 3140 ns
# VSIM: Simulation has finished. There are no more test vectors to simulate.
# VSIM: Simulation has finished.
Done
```

Figure 2: 2-2-1 multilayer perceptron with sigmoid activation function.

```
# KERNEL: Kernel process initialization done.
# Allocation: Simulator allocated 4685 kB (elbread=427 elab2=4122 kernel=135 sdf=0)
# KERNEL: ASDB file was created in location /home/runner/dataset.asdb
# KERNEL: Test Result:
# KERNEL: Inputs: x1 = 0.500000, x2 = 0.500000
# KERNEL: Expected output: 0.000000
# KERNEL: Actual output: 130.871094
# RUNTIME: Info: RUNTIME_0068 testbench.sv (125): $finish called.
# KERNEL: Time: 1140 ns, Iteration: 0, Instance: /mlp_tb, Process: @INITIAL#40_1@.
# KERNEL: stopped at time: 1140 ns
# VSIM: Simulation has finished. There are no more test vectors to simulate.
# VSIM: Simulation has finished.
Done
```

Figure 3: 2-2-1 multilayer perceptron with identity activation function.

## 8 Conclusion

This design demonstrates how a simple MLP can be realized in RTL with fixed-point arithmetic and lookup-table activations. It highlights the mapping of continuous gradient-descent formulas into discrete clocked updates, suitable for FPGA or ASIC prototyping.

## 9 Link to Code

### 9.1 Implementation Variants

There are two EDA Playground versions of the 2-2-1 MLP:

- **Identity Activation:**

Uses pure combinational logic for forward pass fewer LUTs than sigmoid version

Faster convergence but requires output thresholding

- **Sigmoid LUT:**

Biologically plausible activation

Built-in output normalization (0-1 range)

Requires 256-entry ROM but enables deeper networks

- 2-2-1 MLP without sigmoid activation function

- 2-2-1 MLP with sigmoid activation function

## References

- [1] “Multilayer perceptron,” Wikipedia, 2025. [https://en.wikipedia.org/wiki/Multilayer\\_perceptron](https://en.wikipedia.org/wiki/Multilayer_perceptron)
- [2] A. Hinton et al., “A tutorial on multilayer perceptrons,” Neural Computation, 2012.
- [3] “Backpropagation,” Wikipedia, 2025. <https://en.wikipedia.org/wiki/Backpropagation>
- [4] Y. LeCun et al., “Gradient-based learning applied to document recognition,” Proc. IEEE, 1998.
- [5] Pathmind, “A Beginner’s Guide to Multilayer Perceptrons,” 2024. <https://wiki.pathmind.com/multilayer-perceptron>  
ChatGPT for basic knowledge acquisition.